

# Using Neural Networks for Time-Series Prediction

---

Joe Jevnik

November 27, 2017

PyData NYC 2017

1. Introduce problem
2. Phrase problem as an ML problem
3. Collect and apply data
4. Select features
5. Train the model
6. Improve the model
7. Tips

# Who am I?




Quantopian

- Works with storage and manipulation of time series data
- Integrates third-party data sets


# The Problem

---






**Figure 1:** A hit circle




**Figure 1:** A hit circle




**Figure 2:** A slider


# Elements



**Figure 1:** A hit circle



**Figure 2:** A slider



**Figure 3:** Mouse cursor



## Scoring

300 points if on the beat

## Scoring

300 points if on the beat

100 points if slightly off the beat

# Osu! (cont.)

## Scoring

300 points if on the beat

100 points if slightly off the beat



# Osu! (cont.)

## Scoring

300 points if on the beat

100 points if slightly off the beat



50 points if really off the beat

# Osu! (cont.)

## Scoring

300 points if on the beat

100 points if slightly off the beat



50 points if really off the beat



# Osu! (cont.)

## Scoring

300 points if on the beat

100 points if slightly off the beat



50 points if really off the beat



0 points for missing entirely

# Osu! (cont.)

## Scoring

300 points if on the beat

100 points if slightly off the beat



50 points if really off the beat



0 points for missing entirely



## Sample

```
$ mpv videos/osu-example.avi
```



# Problem

Improve my rank quickly

# Problem

Improve my rank quickly

Many new songs a week

# Problem

Improve my rank quickly

Many new songs a week

Particular playstyle

# Problem

Improve my rank quickly

Many new songs a week

Particular playstyle

Arbitrage opportunity?

# Phrasing the Problem

---

# Problem

**Predict my score on a beatmap**

# Problem

**Predict my score on a beatmap**

Need to compute accuracy %

**Predict my score on a beatmap**

Need to compute accuracy %

Temporal Accuracy



## **Predict my score on a beatmap**

Need to compute accuracy %

Temporal Accuracy

Aim Accuracy

## **Classifiers**

Label a sample as a member of one of a finite set of classes.

## **Regressors**

Approximate a numerical function.

Order dependent data

Order dependent data

Sequence of observations

Order dependent data

Sequence of observations

Uses windows of time-sorted observations

# Our Problem

For each hit-object, predict..

# Our Problem

For each hit-object, predict..

## **Classifier**

A label.

- 300
- 100
- 50
- 0 (miss)



# Our Problem

For each hit-object, predict..

## **Classifier**

A label.

- 300
- 100
- 50
- 0 (miss)

## **Regressor**

A numeric error metric.

1. Aim Error ((x, y) error)
2. Accuracy Error (punctuality)

# Data Collection

---

Hit objects in (x, y, time) space.

Hit objects in (x, y, time) space.

Circle Size (CS)

Hit objects in (x, y, time) space.

Circle Size (CS)

Approach Rate (AR)

Hit objects in (x, y, time) space.

Circle Size (CS)

Approach Rate (AR)

Overall Difficulty (OD) (score thresholds)

# Raw Data

## Beatmap

[HitObjects]

103,272,52926,6,0,L|111:176,1,67.5000025749208

93,95,53279,1,2,0:3:0:0:

194,131,53455,2,0,B|263:160|264:100|337:135,1,135.000005149842,0

437,204,53985,2,0,L|432:286,1,67.5000025749208

394,105,54338,6,0,L|399:17,1,67.5000025749208

286,62,54690,1,2,0:3:0:0:

177,54,54867,2,0,B|110:74|110:30|41:53,1,135.000005149842,0

70,213,55396,2,0,L|77:132,1,67.5000025749208

161,215,55749,6,0,P|175:273|247:314,1,135.000005149842,0|2

341,286,56279,1,0,0:0:0:0:

308,183,56455,2,0,P|268:201|245:238,1,67.5000025749208,0|2

- Hard Rock (HR)



- Hard Rock (HR)
- Double Time (DT)

- Hard Rock (HR)
- Double Time (DT)
- Hidden (HD)

- Hard Rock (HR)
- Double Time (DT)
- Hidden (HD)
- etc...

**Time series of...**

**Time series of...**

Cursor location

## Time series of...

Cursor location

Keyboard state

I had about seven years of replays laying around!


## Understanding our data

### Accuracy thresholds (milliseconds)

OD	300	100	50
1	73.5	131.5	189.5
2	67.5	123.5	179.5
3	61.5	115.5	169.5
4	55.5	107.5	159.5
5	49.5	99.5	149.5
6	43.5	91.5	139.5
7	37.5	83.5	129.5
8	31.5	75.5	119.5
9	25.5	67.5	109.5
10	19.5	59.5	99.5



## Understanding our data (cont.)



## Joining Data

## Joining Data

Find all clicks by taking times where key state changes

## Joining Data

Find all clicks by taking times where key state changes

Match click with the nearest hit object (ignores hit locking!)

## **Accuracy Error**

Absolute difference in time.

## **Accuracy Error**

Absolute difference in time.

Comparable across different OD

## Aim Error

Euclidean distance between click and center of circle.

## Aim Error

Euclidean distance between click and center of circle.

Comparable across different CS



# Feature Selection

---

# What is a feature?

Numeric inputs to the ML model

# What is a feature?

Numeric inputs to the ML model

What are we observing

# What is a feature?

Numeric inputs to the ML model

What are we observing

Focus the model on aspects of the data

# What is a feature?

Numeric inputs to the ML model

What are we observing

Focus the model on aspects of the data

Chance to use domain knowledge

## Raw Data

### Beatmap

[HitObjects]

103,272,52926,6,0,L|111:176,1,67.5000025749208

93,95,53279,1,2,0:3:0:0:

194,131,53455,2,0,B|263:160|264:100|337:135,1,135.000005149842,0

437,204,53985,2,0,L|432:286,1,67.5000025749208

394,105,54338,6,0,L|399:17,1,67.5000025749208

286,62,54690,1,2,0:3:0:0:

177,54,54867,2,0,B|110:74|110:30|41:53,1,135.000005149842,0

70,213,55396,2,0,L|77:132,1,67.5000025749208

161,215,55749,6,0,P|175:273|247:314,1,135.000005149842,0|2

341,286,56279,1,0,0:0:0:0:

308,183,56455,2,0,P|268:201|245:238,1,67.5000025749208,0|2

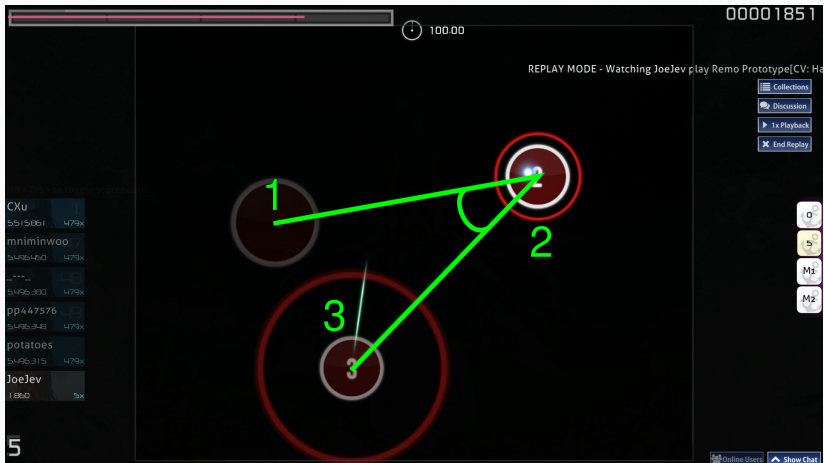
## Simple Features

`absolute_x`

`absolute_y`

`absolute_time`

# Domain Specific Features





time	x	y
00:37.366	372	94
00:37.763	447	205
00:38.027	217	299
00:38.291	229	171
00:38.424	274	358
00:38.688	149	221
00:38.952	330	186
00:39.217	233	127
00:39.481	233	127
00:39.613	198	303

time	x	y
00:37.366	372	94
00:37.763	447	205
00:38.027	217	299
00:38.291	229	171
00:38.424	274	358
00:38.688	149	221
00:38.952	330	186
00:39.217	233	127
00:39.481	233	127
00:39.613	198	303

time	x	y
00:37.366	372	94
00:37.763	447	205
00:38.027	217	299
00:38.291	229	171
00:38.424	274	358
00:38.688	149	221
00:38.952	330	186
00:39.217	233	127
00:39.481	233	127
00:39.613	198	303

time	x	y	relative x	relative y
00:37.366	372	94	-	-
00:37.763	447	205	298	-16
00:38.027	217	299	68	78
00:38.291	229	171	80	-50
00:38.424	274	358	125	137
00:38.688	149	221	0	0
00:38.952	330	186	181	-35
00:39.217	233	127	84	-94
00:39.481	233	127	-	-
00:39.613	198	303	-	-

time	x	y	relative x	relative y
00:37.366	372	94	-	-
00:37.763	447	205	-	-
00:38.027	217	299	-113	113
00:38.291	229	171	-101	-15
00:38.424	274	358	-56	172
00:38.688	149	221	-181	35
00:38.952	330	186	0	0
00:39.217	233	127	-97	-59
00:39.481	233	127	-97	-59
00:39.613	198	303	-	-

time	x	y	relative x	relative y
00:37.366	372	94	-	-
00:37.763	447	205	-	-
00:38.027	217	299	-	-
00:38.291	229	171	-4	44
00:38.424	274	358	41	231
00:38.688	149	221	-84	94
00:38.952	330	186	97	59
00:39.217	233	127	0	0
00:39.481	233	127	0	0
00:39.613	198	303	-35	176

# Osu! Features

- `absolute_x`
- `absolute_y`
- `absolute_time`
- `relative_x`
- `relative_y`
- `relative_time`
- `is_slider_tick`
- `approach_rate`
- `distance_from_previous`
- `distance_to_next`
- `pitch`
- `roll`
- `yaw`

# Training

---



# Input Shapes

**Feature array shape**

( , , , )

# Input Shapes

**Feature array shape**

( , , number of features, )

## Feature array shape

( , window length, number of features, )

## Feature array shape

( number of windows, window length, number of features, )

## Feature array shape

( number of windows, window length, number of features, )

## Label array shape

( number of windows, )

( number of windows, )

```
input_ = keras.layers.Input(  
    shape=(window_length, len(features))  
)
```

```
input_ = keras.layers.Input(  
    shape=(window_length, len(features))  
)  
lstm = keras.layers.LSTM(lstm_layer_size)(input_)
```

```
input_ = keras.layers.Input(  
    shape=(window_length, len(features))  
)  
lstm = keras.layers.LSTM(lstm_layer_size)(input_)  
aim_error = keras.layers.Dense(  
    1,  
    activation='linear',  
    name='aim_error',  
) (lstm)
```



```
input_ = keras.layers.Input(  
    shape=(window_length, len(features))  
)  
lstm = keras.layers.LSTM(lstm_layer_size)(input_)  
aim_error = keras.layers.Dense(  
    1,  
    activation='linear',  
    name='aim_error',  
) (lstm)  
accuracy_error = keras.layers.Dense(  
    1,  
    activation='linear',  
    name='accuracy_error',  
) (lstm)
```

```
model = keras.models.Model(  
    inputs=input_,  
    outputs=[aim_error, accuracy_error],  
)
```

```
model = keras.models.Model(  
    inputs=input_,  
    outputs=[aim_error, accuracy_error],  
)  
  
model.compile(  
    loss='mse',  
    optimizer='rmsprop',  
)
```

```
# compute features, aim_error, accuracy_error
model.fit(
    features,
    {
        'aim_error': aim_error,
        'accuracy_error': accuracy_error,
    },
)
```

```
# compute features, aim_error, accuracy_error
model.fit(
    features,
    {
        'aim_error': aim_error,
        'accuracy_error': accuracy_error,
    },
)
model.predict(features)
```

# How does it look?

## How does it look?

bad

Sensitive to input ranges



Sensitive to input ranges

```
(data - data.mean()) / data.std()
```

# Feature Scaling

Sensitive to input ranges

```
(data - data.mean()) / data.std()
```

Save the mean and std of the training data!

## Data (again)

features:

absolute\_x:

mean: 256.93

std: 581144.54

min: -26.25

max: 42978020236964152.0

absolute\_y:

mean: 188.67

std: 96280.10

min: -12.20

max: 10754663190171874.0

## Data (again)

features:

absolute\_x:

mean: 256.93

std: 581144.54

min: -26.25

max: 42978020236964152.0

absolute\_y:

mean: 188.67

std: 96280.10

min: -12.20


max: 10754663190171874.0

osu! playfield:  
(512, 384)


# Slider Curves

The screenshot displays a music production software interface with a dark theme. At the top, there are four tabs: **compose**, **design**, **timing**, and **song setup**. Below the tabs is a timeline with a sequence of colored circles (blue, green, red) and a section labeled **Break**. On the right side, there is a **Beat Snap Divisor** control set to **1/4** and a button labeled **Insert Break Time**. The main workspace is a grid with four numbered circles (1, 2, 3, 4) connected by lines, representing a sequence of notes or events. The bottom bar shows a timeline with a playhead at **00:23:968** and a **Playback Rate** of **12.0%**. The interface also includes a left sidebar with controls for **Sampleset**, **Auto**, **Additions**, and **Auto**, and a right sidebar with controls for **New**, **Combo**, **Whistle**, **Finish**, **Clap**, **Grid**, **Snap**, **Distance**, **Snap**, **Lock**, and **Notes**.


# Slider Curves



# Slider Curves




# Slider Curves





# Slider Curves




## Understanding our data


### Accuracy thresholds (milliseconds)

OD	300	100	50
1	73.5	131.5	189.5
2	67.5	123.5	179.5
3	61.5	115.5	169.5
4	55.5	107.5	159.5
5	49.5	99.5	149.5
6	43.5	91.5	139.5
7	37.5	83.5	129.5
8	31.5	75.5	119.5
9	25.5	67.5	109.5
10	19.5	59.5	99.5

## Data (again)




## Data (again)




## Fit a distribution

```
lain/lstm.py:605
```



# Sample Weights

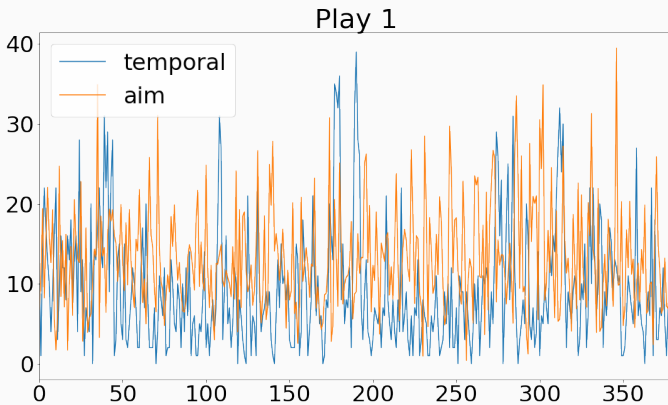


## Sample Weights (cont.)


```
model.fit(  
    ...  
    {  
        'aim_error': aim_error_weights,  
        'accuracy_error': accuracy_error_weights,  
    },  
)
```




## Example Output



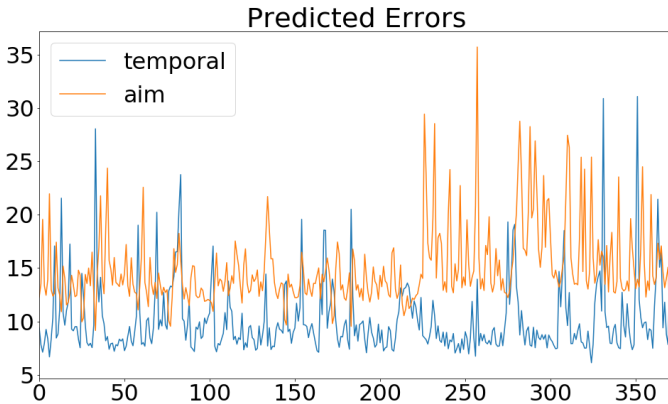
## Example Output



## Example Output



## Example Output



## Example Output (cont.)

Prediction	Actual	Stars
99.85%	99.13%	5.34
99.85%	99.08%	5.34
99.85%	98.51%	5.34
95.95%	96.57%	6.03
98.30%	97.09%	6.24

## Tips

---

verbose flags that log information

# Understanding the process

verbose flags that log information  
progress indicators (`click.progressbar`)



# Understanding the process

verbose flags that log information  
progress indicators (`click.progressbar`)  
print summary statistics early in process

Group code into a domain-aware `Model` class.

Group code into a domain-aware `Model` class.

feature extraction

Group code into a domain-aware `Model` class.

- feature extraction

- feature scaling

Group code into a domain-aware `Model` class.

- feature extraction

- feature scaling

- label extraction

Group code into a domain-aware `Model` class.

- feature extraction

- feature scaling

- label extraction

- managing keras

save models to disk

save models to disk

train on ec2 and use locally



save models to disk

train on ec2 and use locally

train locally then deploy

save models to disk

train on ec2 and use locally

train locally then deploy

supported by keras

# Key Points

1. Most of the work is before or after keras

# Key Points


1. Most of the work is before or after keras
2. Understand the data and the data collection processes

# Key Points

1. Most of the work is before or after keras
2. Understand the data and the data collection processes
3. Osu! is a fun game

# Thank You

Questions?



[github.com/llllllllllll](https://github.com/llllllllllll) (**10 lowercase L's**)

- /lain (model implementation)
- /slider (tools for working with osu! data and API)
- /combine (irc server running lain-as-a-service)